

Run plan (the grid)

Updated 2026-09-20 09:30 from the repo state (results/, config.py, scripts/gpu_box.sh).

Supersedes the earlier draft. Status:  complete ·  partial ·  queued.

One raw prediction per cell (label, probability per label, latency, tokens, paid and list cost) saved to results/<corpus>/<arm>/<model>[__tag].jsonl. All metrics are computed afterward from those files.

Datasets

Code	Set	Docs	Questions	Gold
V	Synthetic: Veridian ApexHip MDL	1,954	10 RFPs (binary; gray flag)	Spec labels from the manifest, relabeled at spec level and text-audited by the mid-tier LLMs
M	Real: Mallinckrodt opioid emails (OIDA)	1,840	4 issues × (broad, narrow) = 8	3-model LLM panel (Sonnet 5, Terra, Gemini 3.8 Flash); Jev and Laya never feed gold

Subsets: V pilot 120 docs; V literal 600; V/M local 400 each (for slow local models); V/M fine-tune splits 30% train / 70% test (586/1,368 and 552/1,288).

Arms: **single** = one question per call; **multi** = all questions for a document in one call.

Models

Group	Models	Where it runs	Setting
Jev	jev-1.13.0 (+ jev-preview)	TypeSafe API	12 ablation variants, see below
Cloud small	Claude Haiku 4.5, GPT-5.6 Luna, Gemini 3.5 Flash-Lite	vendor APIs (batch + cache)	floor effort
Cloud mid	Claude Sonnet 5, GPT-5.6 Terra, Gemini 3.8 Flash	vendor APIs (batch + cache)	floor effort
Laya (System-One competitor)	laya, laya-typed, laya-multilingual	local Mac / rented A100	11 zero-shot variants
Laya fine-tuned (supervised)	laya-ft-veridian, laya-ft-mnk	A100	trained on 30% split, scored on 70%
Local open LLM	Gemma 3 12B (Ollama)	Mac / A100	temp 0, no thinking
Baseline	lexical (keyword overlap, no model)	local	—

Cloud LLM runs

Run	Set	Models	Arm	Effort	Cells / model	Purpose	Status
L1	V	6 cloud	single	floor	19,540	Headline LLM quality	
L2	V	6 cloud	multi	floor	1,954	Batching effect	
L3	M	6 cloud	single	floor	14,720	Headline on real data	
L4	M	6 cloud	multi	floor	1,840	Batching effect, real data	

Run	Set	Models	Arm	Effort	Cells / model	Purpose	Status
L5	V pilot (120)	6 cloud	single	floor vs vendor default	1,200 each	Does default effort beat floor?	✓
L6	V literal (600)	3 small	single	floor	6,000	Jev's "literal" phrasing applied to LLMs (fairness)	✓
G1	V (all)	3 mid	multi	floor	1,954	Gold: spec-level relabel + text audit	✓
G2	M (all)	3 mid	multi	floor	1,840	Gold: Mallinckrodt panel	✓

Jev runs (every variant, both sets, both arms)

Baseline then one lever changed at a time. base / choice / score are on every document; the other levers ran on subsets first and are being filled in to full coverage.

Variant	Lever	Setting	V single	V multi	M single	M multi
jev@base	—	Noul; verbatim RFP; label descriptions as criteria; structured state w/ context; jev-1.13.0	✓	✓	✓	✓
jev@choice	form	Choice	✓	✓	✓	✓
jev@score	form	Score, 5-level rubric	✓	✓	✓	✓
jev@crit_none	criteria	none	🕒	🕒	🕒	🕒
jev@crit_struct	criteria	structured JSON	🕒	🕒	🕒	🕒
jev@literal	phrasing	literal condition w/ boundary cases	🕒	🕒	🕒	🕒
jev@no_context	state	matter context off	🕒	🕒	🕒	🕒
jev@state_string	state	plain string	🕒	🕒	🕒	🕒
jev@gate	composition	relevance-gate Noul	🕒	🕒	🕒	🕒
jev@ensemble	composition	3 phrasings averaged	🕒	🕒	🕒	🕒
jev@decompose	composition	subpart Nouls OR'd in code	🕒	🕒	🕒	🕒
jev@preview	version	jev-preview alias	🕒	🕒	🕒	🕒
jev@recipe	combined	best of the above, fixed after reading the V table	🕒	🕒	🕒	🕒

(Headers on/off is not a Jev lever: M email headers are part of the document text for every model.)

Laya runs (local; \$0 API)

15 variants × 2 arms × 2 corpora, all documents. Laya's 512-token window truncates most instructions and ~half the documents, so two Laya-only levers exist: compact (short instructions/criteria) and chunk (sliding window, max-pool).

Variant	Lever	Status
laya@base, @choice, @score, @literal, @gate, @ensemble, @decompose	same levers as Jev	🕒 (V single base/compact and V multi base done locally; rest queued on A100)

Variant	Lever	Status
laya@compact, @chunk, @recipe (compact+chunk), @recipe_choice	Laya-only window levers	🕒
laya-typed@base, @recipe	typed-decisions checkpoint	🕒
laya-multilingual@base, @recipe	multilingual checkpoint (1024 ctx)	🕒
laya-ft-veridian@compact, @recipe	SUPERVISED fine-tune on V train split, scored on V test split	🕒
laya-ft-mnk@compact, @recipe	SUPERVISED fine-tune on M train split, scored on M test split	🕒
latency sample	laya@base, @recipe, laya-typed@base at concurrency 1 on 400-doc subsets	🕒

Other local runs

Run	Set	Model	Arm	Status
O1	V, M all docs	lexical baseline	single + multi	✅
O2	V local (400)	Gemma 3 12B	multi	✅
O3	V local (400)	Gemma 3 12B	single	🕒 A100
O4	M local (400)	Gemma 3 12B	single + multi	🕒 A100

Prep jobs (not evaluation runs)

Job	Model	Status
Veridian manifest planner	Sonnet 5	✅
Veridian document writer, 1/3 each	Sonnet 5 / Terra / Gemini 3.8 Flash (author recorded per doc)	✅
Mallinckrodt stratified sample via Solr keyword tiers + narrow-issue top-up	none	✅

Analyses from the saved grid (no API calls)

Per question and pooled: precision, recall, F1, macro-F1, elusion, κ , ROC/PR AUC, Brier, ECE, best-F1 threshold, review fraction at 90%/95% recall; with vs without gray docs; panel-unanimous vs majority gold (M); by author family (V); label/probability inconsistency; inter-model agreement; latency p50/p95 (live samples, not batched rows); spend ledger, paid vs list price; per-issue cutoff slider tool (later).